



# Beyond Network Latency

Josh Clark

Distributed Performance Engineer

[https://github.com/je-clark/sf24eu\\_latency](https://github.com/je-clark/sf24eu_latency)



- M.S. in Computer Engineering
  - Focus on networking, protocol design, and security
- Distributed Performance Engineer
  - Using packets to solve complex performance issues

<https://www.jeclark.net>

<https://github.com/je-clark>


[https://github.com/je-clark/sf24eu\\_latency](https://github.com/je-clark/sf24eu_latency)



- **Network Delay** – the time it takes for a packet to travel between client and server, measured as round trip time
- **Application Delay** – the time it takes for an application to respond to a request. This is measured between the application process receiving a message and sending a response to the local server's network stack
- **Server Delay** – the time it takes a server to transfer messages between the NIC and an application process
- **Client Wait Time** - the client's version of server + application delay, usually influenced by user interaction



# Latency in Normal Operation



- Normal internet latency of 30ms simulated on client using tc
- Python script (request\_code.py) deployed on client to simulate a user
- All packets captured via tcpdump on the client



- 0\_normal\_fast\_clear.pcapng
- 1\_normal\_fast\_cipher.pcapng



Δ Conv	Source	Destination	Protocol		
0.005371000	Client	Server	TCP	1448	.....A....
0.000015000	Client	Server	TLSv1.2	789	.....AP...
0.000506000	Server	Client	TCP	0	.....A....
1.002826000	Server	Client	TCP	2896	.....AP...
0.000001000	Server	Client	TCP	2896	.....AP...

Payload = MSS, no PSH | **client REQ**

Payload < MSS, PSH | **end of REQ**

Payload = 0, ACK | **network delay**

Payload > 0 | **server + app delay**

If you only learn one thing from this talk, let it be the  
pattern



- 2\_net\_delay\_clear.pcapng
- 3\_net\_delay\_cipher.pcapng






- 4\_normal\_slow\_clear.pcapng
- 5\_normal\_slow\_cipher.pcapng



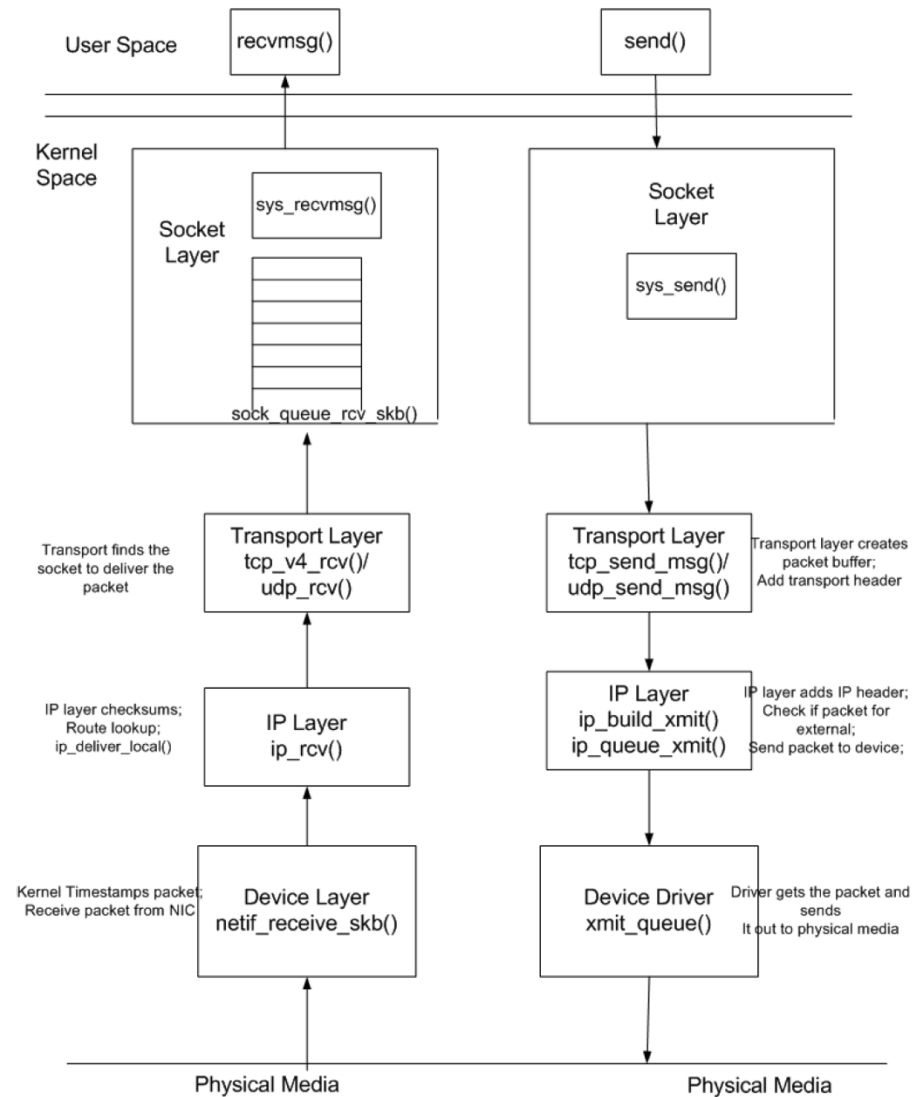
# Latency Under Attack

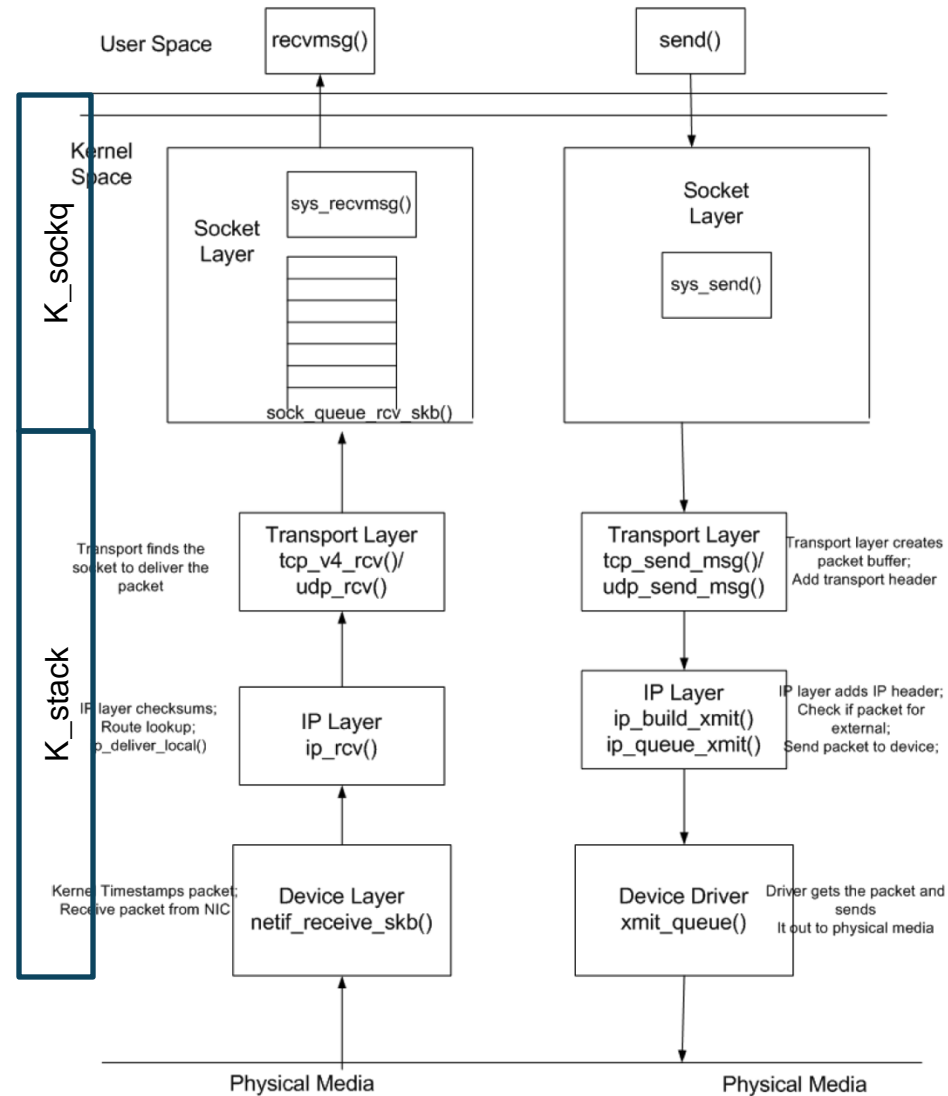


- Switched to Raspberry Pis to reduce overall server capacity
- Switched to Apache because it's less efficient
- Used laptop to add load to the server





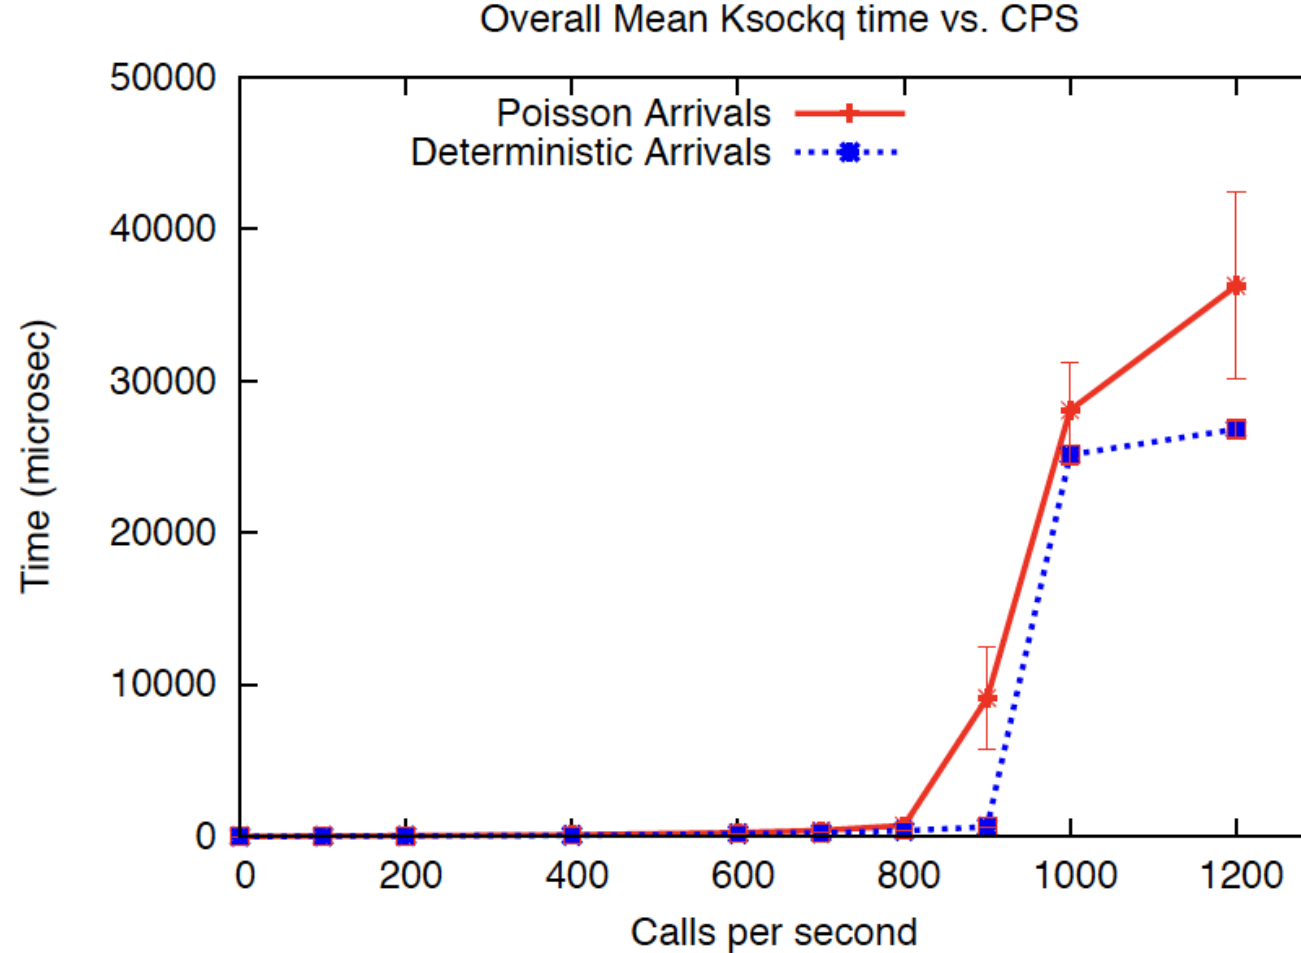
- 6\_get\_flood\_fast\_filtered.pcapng





# Why This Happened





When an application is busy, it can take a long time for the kernel to get the application to pull a message from the socket buffer





- 7\_syn\_flood\_fast\_filtered.pcapng



- To get a packet to the kernel, the NIC must send an interrupt to the CPU
- The process that handles that interrupt is `ksoftirqd`




#b3


PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
14	root	20	0	0	0	0	R	90.7	0.0	0:18.56	ksoftirqd/0
1074	root	20	0	0	0	0	R	9.3	0.0	0:01.33	kworker/0:2+events
709	root	20	0	3116	1840	1420	S	0.7	0.0	0:00.56	dhcpcd
850	root	20	0	0	0	0	I	0.3	0.0	0:00.15	kworker/2:1-events
1069	pi	20	0	9860	3248	2716	R	0.3	0.1	0:05.46	top
1	root	20	0	166824	10172	7436	S	0.0	0.3	0:02.86	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.03	kthreadd
3	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	rcu_gp




# Real World Application



Packet captures from a given layer let us perform latency analysis on both sides of that layer



In a modern internet-facing application, any one of these devices can be causing latency, including devices deployed as appliances. Getting packet captures at multiple layers is critical to isolating a problematic layer.



With a CDN in the mix, it's difficult to determine any latency at the client. Because CDNs operate at Layer 7, we don't even get a good understanding of network latency.

